

For this project, we adopted two different datasets, mentioned below.

ASL alphabet dataset : <https://www.kaggle.com/datasets/grassknoted/asl-alphabet>

Chatbot dataset : <https://www.kaggle.com/datasets/kreeshrajani/3k-conversations-dataset-for-chatbot>

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from collections import Counter
import re
from nltk.corpus import stopwords
stop_words = set(stopwords.words('english'))
import nltk
from nltk.corpus import wordnet
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from datasets import Dataset
from datasets import DatasetDict
import torch
from transformers import AutoModelForSeq2SeqLM, AutoTokenizer, GenerationConfig, TrainingArguments, Trainer
from peft import LoraConfig, get_peft_model, TaskType
from peft import PeftModel
```

```
c:\Users\Paul\anaconda3\envs\torch\lib\site-packages\tqdm\auto.py:21: TqdmWarning: IPProgress not found. Please update
jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

Chatbot Dataset - EDA and preprocessing

```
In [ ]: chatbot_df = pd.read_csv('Conversations_clean.csv')
```

```
In [ ]: chatbot_df.head()
```

Out[]:

	question	answer
0	hi, how are you doing?	i'm fine. how about yourself?
1	i'm fine. how about yourself?	i'm pretty good. thanks for asking.
2	i'm pretty good. thanks for asking.	no problem. so how have you been?
3	no problem. so how have you been?	i've been great. what about you?
4	i've been great. what about you?	i've been good. i'm in school right now.

In []: chatbot_df.shape

Out[]: (3725, 2)

EDA

EDA was performed by following a similar procedure to the one described by Singh (n.d.).

References:





Singh, H. (n.d.). Complete Guide to EDA on Text Data. Kaggle. <https://www.kaggle.com/code/harshsingh2209/complete-guide-to-eda-on-text-data>

```
In [ ]: # Function to count the number of words in a sentence
def count_words(text):
    words = text.split() # extract the words from text
    num_words = len(words) # count the number of words
    return num_words

# Function to count the number of characters in a sentence
def count_characters(text):
    num_char = len(text) # count the number of chatacters
    return num_char

num_words_question = chatbot_df['question'].apply(count_words)
num_words_answer = chatbot_df['answer'].apply(count_words)
num_char_question = chatbot_df['question'].apply(count_characters)
num_char_answer = chatbot_df['answer'].apply(count_characters)
```

```
plt.subplot(2, 2, 1)
sns.boxplot(y = num_words_question)
plt.ylabel('Words per Question')
plt.subplot(2, 2, 2)
sns.boxplot(y = num_words_answer)
plt.ylabel('Words per Answer')
plt.subplot(2, 2, 3)
sns.boxplot(y = num_char_question)
plt.ylabel('Characters per Question')
plt.subplot(2, 2, 4)
sns.boxplot(y = num_char_answer)
plt.ylabel('Characters per Answer')
plt.subplots_adjust(left = 0.1, bottom = 0.1, right = 0.9, top = 0.9, wspace = 0.4, hspace = 0.4)
plt.show()
```



```

In [ ]: # Identify the most frequent words in the question/answer columns

# Function to get the list of words in a sentence
def list_words(text):
    words = text.split() # extract the words from text
    return words


# Function to get the top 20 most common words and their counts
def words_freq(mostcommon):
    words = []
    counts = []
    for word, count in mostcommon:
        words.append(word)
        counts.append(count)
    return words, counts


words_question = chatbot_df['question'].apply(list_words)
words_answer = chatbot_df['answer'].apply(list_words)
corpus_question = []
for jj in range(len(words_question)):
    corpus_question += words_question[jj] # all the words in all the questions
corpus_answer = []
for jj in range(len(words_answer)):
    corpus_answer += words_answer[jj] # all the words in all the answers
mostcommon_words_question = Counter(corpus_question).most_common(20) # 20 most common words in all the questions
mostcommon_words_answer = Counter(corpus_answer).most_common(20) # 20 most common words in all the answers
words_question, counts_question = words_freq(mostcommon_words_question) # top 20 most common words and their counts in questions
words_answer, counts_answer = words_freq(mostcommon_words_answer) # top 20 most common words and their counts in answers

sns.barplot(x = counts_question, y = words_question)
plt.title('Top 20 most common words in questions')
plt.show()

sns.barplot(x = counts_answer, y = words_answer)
plt.title('Top 20 most common words in answers')
plt.show()

```





Preprocessing

In order to replace contractions, we adopted the same procedure described in Replace apostrophe/short words in python (n.d.).

References:

Replace apostrophe/short words in python. (n.d.). Stack Overflow. Retrieved July 24, 2024, from https://owl.purdue.edu/owl/research_and_citation/apa_style/apa_formatting_and_style_guide/reference_list_electronic_sources.html

```
In [ ]: # List of common contractions
contractions = {
    "ain't": "am not / are not",
    "aren't": "are not / am not",
    "can't": "cannot",
```

```
"can't've": "cannot have",
'cause": "because",
"could've": "could have",
"couldn't": "could not",
"couldn't've": "could not have",
"didn't": "did not",
"doesn't": "does not",
"don't": "do not",
"hadn't": "had not",
"hadn't've": "had not have",
"hasn't": "has not",
"haven't": "have not",
"he'd": "he had / he would",
"he'd've": "he would have",
"he'll": "he shall / he will",
"he'll've": "he shall have / he will have",
"he's": "he has / he is",
"how'd": "how did",
"how'd'y": "how do you",
"how'll": "how will",
"how's": "how has / how is",
"i'd": "I had / I would",
"i'd've": "I would have",
"i'll": "I shall / I will",
"i'll've": "I shall have / I will have",
"i'm": "I am",
"i've": "I have",
"isn't": "is not",
"it'd": "it had / it would",
"it'd've": "it would have",
"it'll": "it shall / it will",
"it'll've": "it shall have / it will have",
"it's": "it has / it is",
"let's": "let us",
"ma'am": "madam",
"mayn't": "may not",
"might've": "might have",
"mightn't": "might not",
"mightn't've": "might not have",
"must've": "must have",
"mustn't": "must not",
"mustn't've": "must not have",
```

```
"needn't": "need not",
"needn't've": "need not have",
"o'clock": "of the clock",
"oughtn't": "ought not",
"oughtn't've": "ought not have",
"shan't": "shall not",
"sha'n't": "shall not",
"shan't've": "shall not have",
"she'd": "she had / she would",
"she'd've": "she would have",
"she'll": "she shall / she will",
"she'll've": "she shall have / she will have",
"she's": "she has / she is",
"should've": "should have",
"shouldn't": "should not",
"shouldn't've": "should not have",
"so've": "so have",
"so's": "so as / so is",
"that'd": "that would / that had",
"that'd've": "that would have",
"that's": "that has / that is",
"there'd": "there had / there would",
"there'd've": "there would have",
"there's": "there has / there is",
"they'd": "they had / they would",
"they'd've": "they would have",
"they'll": "they shall / they will",
"they'll've": "they shall have / they will have",
"they're": "they are",
"they've": "they have",
"to've": "to have",
"wasn't": "was not",
"we'd": "we had / we would",
"we'd've": "we would have",
"we'll": "we will",
"we'll've": "we will have",
"we're": "we are",
"we've": "we have",
"weren't": "were not",
"what'll": "what shall / what will",
"what'll've": "what shall have / what will have",
"what're": "what are",
```



```

"what's": "what has / what is",
"what've": "what have",
"when's": "when has / when is",
"when've": "when have",
"where'd": "where did",
"where's": "where has / where is",
"where've": "where have",
"who'll": "who shall / who will",
"who'll've": "who shall have / who will have",
"who's": "who has / who is",
"who've": "who have",
"why's": "why has / why is",
"why've": "why have",
"will've": "will have",
"won't": "will not",
"won't've": "will not have",
"would've": "would have",
"wouldn't": "would not",
"wouldn't've": "would not have",
"y'all": "you all",
"y'all'd": "you all would",
"y'all'd've": "you all would have",
"y'all're": "you all are",
"y'all've": "you all have",
"you'd": "you had / you would",
"you'd've": "you would have",
"you'll": "you shall / you will",
"you'll've": "you shall have / you will have",
"you're": "you are",
"you've": "you have"
}

# Function to replace contractions, remove punctuation and apply lowercase
def clear_text(text):
    for word in text.split(): # remove contractions and apply lowercase
        if word.lower() in contractions:
            text = text.replace(word, contractions[word.lower()])
    text = re.sub(r'^\w\s', '', text) # remove punctuation
    return text

chatbot_df['question'] = chatbot_df['question'].apply(clear_text)

```

```
chatbot_df['answer'] = chatbot_df['answer'].apply(clear_text)
chatbot_df.head()
```

Out []:

	question	answer
0	hi how are you doing	I am fine how about yourself
1	I am fine how about yourself	I am pretty good thanks for asking
2	I am pretty good thanks for asking	no problem so how have you been
3	no problem so how have you been	I have been great what about you
4	I have been great what about you	I have been good I am in school right now

In []:

```
# Tokenization
chatbot_df['token_question'] = chatbot_df['question'].apply(nltk.word_tokenize)
chatbot_df['token_answer'] = chatbot_df['answer'].apply(nltk.word_tokenize)
chatbot_df.head()
```

Out []:

	question	answer	token_question	token_answer
0	hi how are you doing	I am fine how about yourself	[hi, how, are, you, doing]	[I, am, fine, how, about, yourself]
1	I am fine how about yourself	I am pretty good thanks for asking	[I, am, fine, how, about, yourself]	[I, am, pretty, good, thanks, for, asking]
2	I am pretty good thanks for asking	no problem so how have you been	[I, am, pretty, good, thanks, for, asking]	[no, problem, so, how, have, you, been]
3	no problem so how have you been	I have been great what about you	[no, problem, so, how, have, you, been]	[I, have, been, great, what, about, you]
4	I have been great what about you	I have been good I am in school right now	[I, have, been, great, what, about, you]	[I, have, been, good, I, am, in, school, right...]

In []:

```
# Lemmatization

# Function to obtain the right positional tagging prior to Lemmatization
# Same function presented in Python - Lemmatization Approaches with Examples (n.d.)
def right_pos_tagging(tag):
    if tag.startswith('J'):
```

```

        return wordnet.ADJ
    elif tag.startswith('V'):
        return wordnet.VERB
    elif tag.startswith('N'):
        return wordnet.NOUN
    elif tag.startswith('R'):
        return wordnet.ADV
    else:
        return None

# Obtain the right positional tagging prior to Lemmatization
# Positional tags were modified as suggested by Python - Lemmatization Approaches with Examples (n.d.)

tokens = chatbot_df['token_question']
new_tag_tokens = []
for jj in range(len(tokens)):
    tokens_jj = tokens[jj] # tokens at the jjth row
    tag_tokens_jj = nltk.pos_tag(tokens_jj) # POS tags for the generic tokens_jj
    new_tag_tokens.append(list(map(lambda x: (x[0], right_pos_tagging(x[1])), tag_tokens_jj))) # modified POS tags for
tagged_token_question = new_tag_tokens

tokens = chatbot_df['token_answer']
new_tag_tokens = []
for jj in range(len(tokens)):
    tokens_jj = tokens[jj] # tokens at the jjth row
    tag_tokens_jj = nltk.pos_tag(tokens_jj) # POS tags for the generic tokens_jj
    new_tag_tokens.append(list(map(lambda x: (x[0], right_pos_tagging(x[1])), tag_tokens_jj))) # modified POS tags for
tagged_token_answer = new_tag_tokens

# Implement Lemmatization on the tokens
# A procedure similar to the one described in Python - Lemmatization Approaches with Examples (n.d.) and by Kumar (2019)

wnl = WordNetLemmatizer()
Lemmatization = []
for jj in range(len(tagged_token_question)):
    lemmatized_question = []
    # Same lines of codes used in Python - Lemmatization Approaches with Examples (n.d.)
    for word, tag in tagged_token_question[jj]:
        if tag is None:
            lemmatized_question.append(word)
        else:
            lemmatized_question.append(wnl.lemmatize(word, tag))

```

```

    Lemmatization.append(lemmatized_question)
chatbot_df['lem_question'] = Lemmatization

Lemmatization = []
for jj in range(len(tagged_token_answer)):
    lemmatized_answer = []
    # Same lines of codes used in Python - Lemmatization Approaches with Examples (n.d.)
    for word, tag in tagged_token_answer[jj]:
        if tag is None:
            lemmatized_answer.append(word)
        else:
            lemmatized_answer.append(wnl.lemmatize(word, tag))
    Lemmatization.append(lemmatized_answer)
chatbot_df['lem_answer'] = Lemmatization

chatbot_df.head()

# References
# Kumar, R. (2021, August 6). Natural Language Processing | Text Preprocessing | Spacy vs NLTK. Medium. https://medium.com/@rkumar1998/natural-language-processing-text-preprocessing-spacy-vs-nltk-1234567890
# Python - Lemmatization Approaches with Examples. (n.d.). Geeks for Geeks. https://www.geeksforgeeks.org/python-lemmatization-approaches-with-examples/

```

Out[]:

	question	answer	token_question	token_answer	lem_question	lem_answer
0	hi how are you doing	I am fine how about yourself	[hi, how, are, you, doing]	[I, am, fine, how, about, yourself]	[hi, how, be, you, do]	[I, be, fine, how, about, yourself]
1	I am fine how about yourself	I am pretty good thanks for asking	[I, am, fine, how, about, yourself]	[I, am, pretty, good, thanks, for, asking]	[I, be, fine, how, about, yourself]	[I, be, pretty, good, thanks, for, ask]
2	I am pretty good thanks for asking	no problem so how have you been	[I, am, pretty, good, thanks, for, asking]	[no, problem, so, how, have, you, been]	[I, be, pretty, good, thanks, for, ask]	[no, problem, so, how, have, you, be]
3	no problem so how have you been	I have been great what about you	[no, problem, so, how, have, you, been]	[I, have, been, great, what, about, you]	[no, problem, so, how, have, you, be]	[I, have, be, great, what, about, you]
4	I have been great what about you	I have been good I am in school right now	[I, have, been, great, what, about, you]	[I, have, been, good, I, am, in, school, right...]	[I, have, be, great, what, about, you]	[I, have, be, good, I, be, in, school, right, ...]

T5 Fine Tuning

This section was created based on a combination of the following references.

References: <https://www.kaggle.com/code/ajinkyabhandare2002/fine-tune-flan-t5-base-for-chat-with-peft-lora>

Import tokenizer and model

```
In [ ]: model_name = 'google/flan-t5-base'
model = AutoModelForSeq2SeqLM.from_pretrained(model_name, torch_dtype = torch.bfloat16)
tokenizer = AutoTokenizer.from_pretrained(model_name)

In [ ]: # Check number of parameters to train
def model_parameters(model):
    trainable_model_params = 0
    all_model_params = 0
    for _, param in model.named_parameters():
        all_model_params += param.numel()
        if param.requires_grad:
            trainable_model_params += param.numel()
    return f"trainable model parameters: {trainable_model_params}\nall model parameters: {all_model_params}\npercentage of trainable model parameters: {trainable_model_params/all_model_params*100:.2f}%"

print(model_parameters(model))
```

```
trainable model parameters: 247577856
all model parameters: 247577856
percentage of trainable model parameters: 100.00%
```

Preprocess Data for Retrain

```
In [ ]: chatbot_df=chatbot_df.drop(columns=['token_question', 'token_answer', 'lem_question', 'lem_answer'])

In [ ]: train_data, temp_data = train_test_split(chatbot_df, test_size=0.2, random_state=42)
val_data, test_data = train_test_split(temp_data, test_size=0.5, random_state=42)
train_dataset = Dataset.from_pandas(train_data)
val_dataset = Dataset.from_pandas(val_data)
test_dataset = Dataset.from_pandas(test_data)
```

```
In [ ]: from datasets import DatasetDict
working_dataset = DatasetDict({
    "train": train_dataset,
    "validation": val_dataset,
    "test": test_dataset,
})
```

```
In [ ]: working_dataset
```

```
Out[ ]: DatasetDict({
  train: Dataset({
    features: ['question', 'answer', '__index_level_0__'],
    num_rows: 2980
  })
  validation: Dataset({
    features: ['question', 'answer', '__index_level_0__'],
    num_rows: 372
  })
  test: Dataset({
    features: ['question', 'answer', '__index_level_0__'],
    num_rows: 373
  })
})
```

```
In [ ]: working_dataset["train"] = working_dataset["train"].remove_columns("__index_level_0__")
working_dataset["validation"] = working_dataset["validation"].remove_columns("__index_level_0__")
working_dataset["test"] = working_dataset["test"].remove_columns("__index_level_0__")
working_dataset
```

```
Out[ ]: DatasetDict({
  train: Dataset({
    features: ['question', 'answer'],
    num_rows: 2980
  })
  validation: Dataset({
    features: ['question', 'answer'],
    num_rows: 372
  })
  test: Dataset({
    features: ['question', 'answer'],
    num_rows: 373
  })
})
```

```
In [ ]: def tokenize_function(example):
    # start_prompt = 'Answer the following question.\n\n'
    # end_prompt = '\n\nSummary: '
    # prompt = [start_prompt + question + end_prompt for question in example["question"]]
    example['input_ids'] = tokenizer(example['question'], padding='max_length', truncation=True, return_tensors="pt")
    example['labels'] = tokenizer(example["answer"], padding='max_length', truncation=True, return_tensors="pt").input_ids

    return example

# The dataset actually contains 3 diff splits: train, validation, test.
# The tokenize_function code is handling all data across all splits in batches.
tokenized_datasets = working_dataset.map(tokenize_function, batched=True)
tokenized_datasets = tokenized_datasets.remove_columns(['question', 'answer'])
```

```
Map: 100%|██████████| 2980/2980 [00:00<00:00, 5327.93 examples/s]
Map: 100%|██████████| 372/372 [00:00<00:00, 6009.08 examples/s]
Map: 100%|██████████| 373/373 [00:00<00:00, 6893.48 examples/s]
```

```
In [ ]: tokenized_datasets
```

```
Out[ ]: DatasetDict({
  train: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 2980
  })
  validation: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 372
  })
  test: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 373
  })
})
```

```
In [ ]: print(f"Shapes of the datasets:")
        print(f"Dataset: {tokenized_datasets.shape}")
        print(tokenized_datasets)
```

Shapes of the datasets:

Dataset: {'train': (2980, 2), 'validation': (372, 2), 'test': (373, 2)}

```
DatasetDict({
  train: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 2980
  })
  validation: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 372
  })
  test: Dataset({
    features: ['input_ids', 'labels'],
    num_rows: 373
  })
})
```

Setup the PEFT/LoRA model for Fine-Tuning

```
In [ ]: lora_config = LoraConfig(
        r=8, # Rank
        lora_alpha=8,
```



```

target_modules=["q", "v"],
lora_dropout=0.05,
bias="none",
task_type=TaskType.SEQ_2_SEQ_LM # FLAN-T5
)

```

Add LoRA adapter layers/prameters to the LLM model to be trained

```

In [ ]: peft_model = get_peft_model(model, lora_config)
        print(model_parameters(peft_model))

```

```

bin c:\Users\Paul\anaconda3\envs\torch\lib\site-packages\bitsandbytes\libbitsandbytes_cuda118.dll
trainable model parameters: 884736
all model parameters: 248462592
percentage of trainable model parameters: 0.36%

```

Train PEFT Adapter

```

In [ ]: # Define Trianing Arguements
        output_dir = f'./peft-conversation-training'

        peft_training_args = TrainingArguments(
            output_dir=output_dir,
            auto_find_batch_size=True,
            learning_rate=1e-3, # Higher Learning rate than full fine-tuning.
            num_train_epochs=5,
            save_steps=100,
            save_strategy='steps',
            evaluation_strategy='steps',
            eval_steps=10,
        )

        peft_trainer = Trainer(
            model=peft_model,
            args=peft_training_args,
            train_dataset=tokenized_datasets["train"],
            eval_dataset = tokenized_datasets['validation']
        )

```

```
c:\Users\Paul\anaconda3\envs\torch\lib\site-packages\transformers\training_args.py:1525: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of 🤗 Transformers. Use `eval_strategy` instead
warnings.warn(
```

Train the model

```
In [ ]: try:
        peft_trainer.train()
    except OutOfMemoryError:
        print("Training interrupted due to OOM. Saving model checkpoint...")
        peft_model_path="./peft-conversation-checkpoint-local"
        peft_trainer.model.save_pretrained(peft_model_path)
        tokenizer.save_pretrained(peft_model_path)
        print("Checkpoint saved. You can resume training from here.")
```

```
1%|          | 10/1865 [00:02<05:55, 5.21it/s]
```

```
1%|          | 11/1865 [00:05<35:46, 1.16s/it]
```

```
{'eval_loss': 30.860214233398438, 'eval_runtime': 3.1556, 'eval_samples_per_second': 117.886, 'eval_steps_per_second': 14.894, 'epoch': 0.03}
```

```
1%|          | 21/1865 [00:11<39:59, 1.30s/it]
```

```
{'eval_loss': 14.452284812927246, 'eval_runtime': 3.6103, 'eval_samples_per_second': 103.039, 'eval_steps_per_second': 13.018, 'epoch': 0.05}
```

```
2%||         | 31/1865 [00:16<35:48, 1.17s/it]
```

```
{'eval_loss': 4.510080814361572, 'eval_runtime': 3.1676, 'eval_samples_per_second': 117.44, 'eval_steps_per_second': 14.838, 'epoch': 0.08}
```

```
2%||         | 41/1865 [00:21<35:20, 1.16s/it]
```

```
{'eval_loss': 4.096774101257324, 'eval_runtime': 3.1505, 'eval_samples_per_second': 118.075, 'eval_steps_per_second': 14.918, 'epoch': 0.11}
```

```
3%||         | 51/1865 [00:26<35:10, 1.16s/it]
```

```
{'eval_loss': 2.7189180850982666, 'eval_runtime': 3.1536, 'eval_samples_per_second': 117.962, 'eval_steps_per_second': 14.904, 'epoch': 0.13}
```

```
3%||         | 61/1865 [00:31<38:45, 1.29s/it]
```

```
{'eval_loss': 1.0347782373428345, 'eval_runtime': 3.5698, 'eval_samples_per_second': 104.208, 'eval_steps_per_second': 13.166, 'epoch': 0.16}
```

4%|█| 71/1865 [00:36<34:48, 1.16s/it]

{'eval_loss': 0.5599378347396851, 'eval_runtime': 3.1425, 'eval_samples_per_second': 118.377, 'eval_steps_per_second': 14.956, 'epoch': 0.19}

4%|█| 81/1865 [00:41<34:31, 1.16s/it]

{'eval_loss': 0.33358535170555115, 'eval_runtime': 3.1468, 'eval_samples_per_second': 118.214, 'eval_steps_per_second': 14.936, 'epoch': 0.21}

5%|█| 91/1865 [00:46<34:17, 1.16s/it]

{'eval_loss': 0.2514490783214569, 'eval_runtime': 3.1435, 'eval_samples_per_second': 118.34, 'eval_steps_per_second': 14.952, 'epoch': 0.24}

5%|█| 100/1865 [00:51<06:41, 4.40it/s]

{'eval_loss': 0.1840347796678543, 'eval_runtime': 3.1439, 'eval_samples_per_second': 118.326, 'eval_steps_per_second': 14.95, 'epoch': 0.27}

6%|█| 111/1865 [00:57<34:00, 1.16s/it]

{'eval_loss': 0.16665616631507874, 'eval_runtime': 3.1478, 'eval_samples_per_second': 118.177, 'eval_steps_per_second': 14.931, 'epoch': 0.29}

6%|█| 121/1865 [01:02<33:37, 1.16s/it]

{'eval_loss': 0.1538243442773819, 'eval_runtime': 3.1331, 'eval_samples_per_second': 118.731, 'eval_steps_per_second': 15.001, 'epoch': 0.32}

7%|█| 131/1865 [01:07<33:29, 1.16s/it]

{'eval_loss': 0.14440524578094482, 'eval_runtime': 3.1422, 'eval_samples_per_second': 118.39, 'eval_steps_per_second': 14.958, 'epoch': 0.35}

8%|█| 141/1865 [01:12<33:15, 1.16s/it]

{'eval_loss': 0.13718077540397644, 'eval_runtime': 3.1381, 'eval_samples_per_second': 118.544, 'eval_steps_per_second': 14.977, 'epoch': 0.38}

8%|█| 151/1865 [01:17<33:01, 1.16s/it]

{'eval_loss': 0.13080687820911407, 'eval_runtime': 3.1309, 'eval_samples_per_second': 118.817, 'eval_steps_per_second': 15.012, 'epoch': 0.4}

9%|█| 161/1865 [01:22<32:56, 1.16s/it]

{'eval_loss': 0.1274886578321457, 'eval_runtime': 3.1428, 'eval_samples_per_second': 118.367, 'eval_steps_per_second': 14.955, 'epoch': 0.43}

9%|█| 171/1865 [01:27<32:44, 1.16s/it]

{'eval_loss': 0.12263734638690948, 'eval_runtime': 3.1424, 'eval_samples_per_second': 118.381, 'eval_steps_per_second': 14.957, 'epoch': 0.46}

10%|█| 180/1865 [01:32<06:22, 4.40it/s]

{'eval_loss': 0.11838982999324799, 'eval_runtime': 3.5633, 'eval_samples_per_second': 104.399, 'eval_steps_per_second': 13.19, 'epoch': 0.48}

10%|█| 191/1865 [01:37<32:24, 1.16s/it]

{'eval_loss': 0.11515036970376968, 'eval_runtime': 3.1372, 'eval_samples_per_second': 118.578, 'eval_steps_per_second': 14.982, 'epoch': 0.51}

11%|█| 200/1865 [01:42<06:17, 4.41it/s]

{'eval_loss': 0.11255670338869095, 'eval_runtime': 3.1597, 'eval_samples_per_second': 117.734, 'eval_steps_per_second': 14.875, 'epoch': 0.54}

11%|█| 210/1865 [01:47<06:20, 4.35it/s]

{'eval_loss': 0.10888146609067917, 'eval_runtime': 3.1586, 'eval_samples_per_second': 117.772, 'eval_steps_per_second': 14.88, 'epoch': 0.56}

12%|█| 221/1865 [01:53<31:43, 1.16s/it]

{'eval_loss': 0.10646631568670273, 'eval_runtime': 3.1392, 'eval_samples_per_second': 118.5, 'eval_steps_per_second': 14.972, 'epoch': 0.59}

12%|█| 231/1865 [01:58<31:26, 1.15s/it]

{'eval_loss': 0.1039881557226181, 'eval_runtime': 3.1258, 'eval_samples_per_second': 119.008, 'eval_steps_per_second': 15.036, 'epoch': 0.62}

13%|█| 241/1865 [02:03<31:19, 1.16s/it]

{'eval_loss': 0.10217154026031494, 'eval_runtime': 3.1362, 'eval_samples_per_second': 118.615, 'eval_steps_per_second': 14.986, 'epoch': 0.64}

13%|█| 251/1865 [02:08<31:09, 1.16s/it]

{'eval_loss': 0.0994623675942421, 'eval_runtime': 3.1405, 'eval_samples_per_second': 118.452, 'eval_steps_per_second': 14.966, 'epoch': 0.67}

14%|███ | 261/1865 [02:13<31:00, 1.16s/it]
{'eval_loss': 0.09701045602560043, 'eval_runtime': 3.1432, 'eval_samples_per_second': 118.35, 'eval_steps_per_second': 14.953, 'epoch': 0.7}

15%|███ | 271/1865 [02:18<30:54, 1.16s/it]
{'eval_loss': 0.09564012289047241, 'eval_runtime': 3.1535, 'eval_samples_per_second': 117.965, 'eval_steps_per_second': 14.904, 'epoch': 0.72}

15%|███ | 281/1865 [02:23<30:37, 1.16s/it]
{'eval_loss': 0.09327221661806107, 'eval_runtime': 3.1428, 'eval_samples_per_second': 118.366, 'eval_steps_per_second': 14.955, 'epoch': 0.75}

16%|███ | 291/1865 [02:28<30:27, 1.16s/it]
{'eval_loss': 0.09321446716785431, 'eval_runtime': 3.1481, 'eval_samples_per_second': 118.166, 'eval_steps_per_second': 14.93, 'epoch': 0.78}

16%|███ | 300/1865 [02:33<05:55, 4.40it/s]
{'eval_loss': 0.09120358526706696, 'eval_runtime': 3.1441, 'eval_samples_per_second': 118.315, 'eval_steps_per_second': 14.948, 'epoch': 0.8}

17%|███ | 311/1865 [02:38<30:07, 1.16s/it]
{'eval_loss': 0.08915070444345474, 'eval_runtime': 3.148, 'eval_samples_per_second': 118.169, 'eval_steps_per_second': 14.93, 'epoch': 0.83}

17%|███ | 321/1865 [02:43<30:01, 1.17s/it]
{'eval_loss': 0.08911395072937012, 'eval_runtime': 3.1618, 'eval_samples_per_second': 117.656, 'eval_steps_per_second': 14.865, 'epoch': 0.86}

18%|███ | 331/1865 [02:48<29:43, 1.16s/it]
{'eval_loss': 0.08748109638690948, 'eval_runtime': 3.1512, 'eval_samples_per_second': 118.052, 'eval_steps_per_second': 14.915, 'epoch': 0.88}

18%|███ | 341/1865 [02:53<29:30, 1.16s/it]
{'eval_loss': 0.0862787738442421, 'eval_runtime': 3.1497, 'eval_samples_per_second': 118.105, 'eval_steps_per_second': 14.922, 'epoch': 0.91}

19%|███ | 350/1865 [02:58<05:45, 4.39it/s]

{'eval_loss': 0.08584299683570862, 'eval_runtime': 3.1394, 'eval_samples_per_second': 118.494, 'eval_steps_per_second': 14.971, 'epoch': 0.94}

19%|██████████| 361/1865 [03:03<29:02, 1.16s/it]
{'eval_loss': 0.08338584005832672, 'eval_runtime': 3.139, 'eval_samples_per_second': 118.511, 'eval_steps_per_second': 14.973, 'epoch': 0.97}

20%|██████████| 370/1865 [03:08<05:39, 4.40it/s]
{'eval_loss': 0.08290805667638779, 'eval_runtime': 3.3862, 'eval_samples_per_second': 109.857, 'eval_steps_per_second': 13.88, 'epoch': 0.99}

20%|██████████| 381/1865 [03:14<31:50, 1.29s/it]
{'eval_loss': 0.08205749839544296, 'eval_runtime': 3.5065, 'eval_samples_per_second': 106.089, 'eval_steps_per_second': 13.404, 'epoch': 1.02}

21%|██████████| 391/1865 [03:19<28:32, 1.16s/it]
{'eval_loss': 0.08042464405298233, 'eval_runtime': 3.1366, 'eval_samples_per_second': 118.6, 'eval_steps_per_second': 14.984, 'epoch': 1.05}

21%|██████████| 400/1865 [03:24<05:31, 4.42it/s]
{'eval_loss': 0.07961084693670273, 'eval_runtime': 3.1192, 'eval_samples_per_second': 119.262, 'eval_steps_per_second': 15.068, 'epoch': 1.07}

22%|██████████| 410/1865 [03:29<05:32, 4.37it/s]
{'eval_loss': 0.0802408829331398, 'eval_runtime': 3.0914, 'eval_samples_per_second': 120.335, 'eval_steps_per_second': 15.204, 'epoch': 1.1}

23%|██████████| 421/1865 [03:35<27:45, 1.15s/it]
{'eval_loss': 0.07775222510099411, 'eval_runtime': 3.1211, 'eval_samples_per_second': 119.19, 'eval_steps_per_second': 15.059, 'epoch': 1.13}

23%|██████████| 431/1865 [03:40<27:22, 1.15s/it]
{'eval_loss': 0.07681766897439957, 'eval_runtime': 3.1029, 'eval_samples_per_second': 119.887, 'eval_steps_per_second': 15.147, 'epoch': 1.15}

24%|██████████| 441/1865 [03:45<27:32, 1.16s/it]
{'eval_loss': 0.0776052176952362, 'eval_runtime': 3.154, 'eval_samples_per_second': 117.944, 'eval_steps_per_second': 14.902, 'epoch': 1.18}

24%|██████ | 451/1865 [03:50<27:07, 1.15s/it]
{'eval_loss': 0.07609836757183075, 'eval_runtime': 3.1235, 'eval_samples_per_second': 119.098, 'eval_steps_per_second': 15.047, 'epoch': 1.21}

25%|██████ | 461/1865 [03:55<27:02, 1.16s/it]
{'eval_loss': 0.07625062763690948, 'eval_runtime': 3.1318, 'eval_samples_per_second': 118.781, 'eval_steps_per_second': 15.007, 'epoch': 1.23}

25%|██████ | 471/1865 [04:00<26:41, 1.15s/it]
{'eval_loss': 0.07512180507183075, 'eval_runtime': 3.1189, 'eval_samples_per_second': 119.271, 'eval_steps_per_second': 15.069, 'epoch': 1.26}

26%|██████ | 481/1865 [04:05<26:15, 1.14s/it]
{'eval_loss': 0.07464402914047241, 'eval_runtime': 3.0789, 'eval_samples_per_second': 120.824, 'eval_steps_per_second': 15.265, 'epoch': 1.29}

26%|██████ | 491/1865 [04:10<26:20, 1.15s/it]
{'eval_loss': 0.07493279874324799, 'eval_runtime': 3.136, 'eval_samples_per_second': 118.624, 'eval_steps_per_second': 14.987, 'epoch': 1.31}

27%|██████ | 500/1865 [04:11<04:59, 4.56it/s]
{'loss': 1.6373, 'grad_norm': 0.07249174267053604, 'learning_rate': 0.0007319034852546918, 'epoch': 1.34}

27%|██████ | 500/1865 [04:14<04:59, 4.56it/s]
{'eval_loss': 0.0734049454331398, 'eval_runtime': 3.1305, 'eval_samples_per_second': 118.831, 'eval_steps_per_second': 15.014, 'epoch': 1.34}

27%|██████ | 511/1865 [04:20<26:12, 1.16s/it]
{'eval_loss': 0.07353095710277557, 'eval_runtime': 3.1432, 'eval_samples_per_second': 118.352, 'eval_steps_per_second': 14.953, 'epoch': 1.37}

28%|██████ | 521/1865 [04:25<26:37, 1.19s/it]
{'eval_loss': 0.07254914194345474, 'eval_runtime': 3.2222, 'eval_samples_per_second': 115.449, 'eval_steps_per_second': 14.586, 'epoch': 1.39}

28%|██████ | 531/1865 [04:30<26:04, 1.17s/it]
{'eval_loss': 0.07197685539722443, 'eval_runtime': 3.1853, 'eval_samples_per_second': 116.786, 'eval_steps_per_second': 14.755, 'epoch': 1.42}

29%|██████ | 541/1865 [04:35<25:45, 1.17s/it]

{'eval_loss': 0.0715988352894783, 'eval_runtime': 3.1645, 'eval_samples_per_second': 117.553, 'eval_steps_per_second': 14.852, 'epoch': 1.45}

30%|██████ | 551/1865 [04:40<25:11, 1.15s/it]

{'eval_loss': 0.07134681940078735, 'eval_runtime': 3.1136, 'eval_samples_per_second': 119.474, 'eval_steps_per_second': 15.095, 'epoch': 1.47}

30%|██████ | 561/1865 [04:45<24:54, 1.15s/it]

{'eval_loss': 0.0711420550942421, 'eval_runtime': 3.1238, 'eval_samples_per_second': 119.087, 'eval_steps_per_second': 15.046, 'epoch': 1.5}

31%|██████ | 571/1865 [04:50<24:46, 1.15s/it]

{'eval_loss': 0.07077453285455704, 'eval_runtime': 3.1249, 'eval_samples_per_second': 119.045, 'eval_steps_per_second': 15.041, 'epoch': 1.53}

31%|██████ | 581/1865 [04:55<24:46, 1.16s/it]

{'eval_loss': 0.06965620815753937, 'eval_runtime': 3.155, 'eval_samples_per_second': 117.908, 'eval_steps_per_second': 14.897, 'epoch': 1.55}

32%|██████ | 591/1865 [05:00<24:31, 1.15s/it]

{'eval_loss': 0.06988197565078735, 'eval_runtime': 3.1339, 'eval_samples_per_second': 118.703, 'eval_steps_per_second': 14.997, 'epoch': 1.58}

32%|██████ | 600/1865 [05:05<04:45, 4.42it/s]

{'eval_loss': 0.06925718486309052, 'eval_runtime': 3.1055, 'eval_samples_per_second': 119.786, 'eval_steps_per_second': 15.134, 'epoch': 1.61}

33%|██████ | 611/1865 [05:10<24:07, 1.15s/it]

{'eval_loss': 0.06922043114900589, 'eval_runtime': 3.1206, 'eval_samples_per_second': 119.208, 'eval_steps_per_second': 15.061, 'epoch': 1.64}

33%|██████ | 621/1865 [05:15<23:36, 1.14s/it]

{'eval_loss': 0.06869539618492126, 'eval_runtime': 3.0973, 'eval_samples_per_second': 120.106, 'eval_steps_per_second': 15.175, 'epoch': 1.66}

34%|██████ | 631/1865 [05:20<23:35, 1.15s/it]

{'eval_loss': 0.0683436244726181, 'eval_runtime': 3.1211, 'eval_samples_per_second': 119.187, 'eval_steps_per_second': 15.059, 'epoch': 1.69}

34%|███████| 641/1865 [05:25<23:12, 1.14s/it]

{'eval_loss': 0.06796559691429138, 'eval_runtime': 3.0717, 'eval_samples_per_second': 121.106, 'eval_steps_per_second': 15.301, 'epoch': 1.72}

35%|███████| 651/1865 [05:30<23:02, 1.14s/it]

{'eval_loss': 0.06760332733392715, 'eval_runtime': 3.0862, 'eval_samples_per_second': 120.535, 'eval_steps_per_second': 15.229, 'epoch': 1.74}

35%|███████| 661/1865 [05:35<22:51, 1.14s/it]

{'eval_loss': 0.06821761280298233, 'eval_runtime': 3.0976, 'eval_samples_per_second': 120.092, 'eval_steps_per_second': 15.173, 'epoch': 1.77}

36%|███████| 671/1865 [05:40<22:50, 1.15s/it]

{'eval_loss': 0.06681577861309052, 'eval_runtime': 3.1167, 'eval_samples_per_second': 119.359, 'eval_steps_per_second': 15.08, 'epoch': 1.8}

37%|███████| 681/1865 [05:45<23:15, 1.18s/it]

{'eval_loss': 0.06684202700853348, 'eval_runtime': 3.2144, 'eval_samples_per_second': 115.73, 'eval_steps_per_second': 14.622, 'epoch': 1.82}

37%|███████| 691/1865 [05:50<22:42, 1.16s/it]

{'eval_loss': 0.06701003760099411, 'eval_runtime': 3.1479, 'eval_samples_per_second': 118.173, 'eval_steps_per_second': 14.93, 'epoch': 1.85}

38%|███████| 700/1865 [05:55<04:23, 4.42it/s]

{'eval_loss': 0.06670551747083664, 'eval_runtime': 3.1369, 'eval_samples_per_second': 118.588, 'eval_steps_per_second': 14.983, 'epoch': 1.88}

38%|███████| 711/1865 [06:00<22:48, 1.19s/it]

{'eval_loss': 0.06666351854801178, 'eval_runtime': 3.2274, 'eval_samples_per_second': 115.262, 'eval_steps_per_second': 14.563, 'epoch': 1.9}

39%|███████| 721/1865 [06:05<22:00, 1.15s/it]

{'eval_loss': 0.06557144969701767, 'eval_runtime': 3.1177, 'eval_samples_per_second': 119.318, 'eval_steps_per_second': 15.075, 'epoch': 1.93}

39%|██████ | 731/1865 [06:10<21:51, 1.16s/it]

{'eval_loss': 0.06646400690078735, 'eval_runtime': 3.1507, 'eval_samples_per_second': 118.07, 'eval_steps_per_second': 14.917, 'epoch': 1.96}

40%|██████ | 741/1865 [06:15<21:38, 1.15s/it]

{'eval_loss': 0.06565020233392715, 'eval_runtime': 3.1346, 'eval_samples_per_second': 118.674, 'eval_steps_per_second': 14.994, 'epoch': 1.98}

40%|██████ | 751/1865 [06:20<21:11, 1.14s/it]

{'eval_loss': 0.06534568220376968, 'eval_runtime': 3.0999, 'eval_samples_per_second': 120.002, 'eval_steps_per_second': 15.162, 'epoch': 2.01}

41%|██████ | 761/1865 [06:25<20:55, 1.14s/it]

{'eval_loss': 0.0649729073047638, 'eval_runtime': 3.0918, 'eval_samples_per_second': 120.319, 'eval_steps_per_second': 15.202, 'epoch': 2.04}

41%|██████ | 771/1865 [06:31<24:36, 1.35s/it]

{'eval_loss': 0.06561870127916336, 'eval_runtime': 3.7772, 'eval_samples_per_second': 98.486, 'eval_steps_per_second': 12.443, 'epoch': 2.06}

42%|██████ | 781/1865 [06:36<21:20, 1.18s/it]

{'eval_loss': 0.06476814299821854, 'eval_runtime': 3.1965, 'eval_samples_per_second': 116.379, 'eval_steps_per_second': 14.704, 'epoch': 2.09}

42%|██████ | 790/1865 [06:41<03:57, 4.52it/s]

{'eval_loss': 0.0649571567773819, 'eval_runtime': 3.1845, 'eval_samples_per_second': 116.815, 'eval_steps_per_second': 14.759, 'epoch': 2.12}

43%|██████ | 800/1865 [06:46<04:09, 4.27it/s]

{'eval_loss': 0.06444787234067917, 'eval_runtime': 3.2269, 'eval_samples_per_second': 115.28, 'eval_steps_per_second': 14.565, 'epoch': 2.14}

43%|██████ | 811/1865 [06:51<20:49, 1.19s/it]

{'eval_loss': 0.06427986174821854, 'eval_runtime': 3.2121, 'eval_samples_per_second': 115.813, 'eval_steps_per_second': 14.632, 'epoch': 2.17}

44%|██████ | 821/1865 [06:56<20:19, 1.17s/it]

```
{'eval_loss': 0.06420110911130905, 'eval_runtime': 3.11, 'eval_samples_per_second': 119.614, 'eval_steps_per_second': 15.113, 'epoch': 2.2}
```

45%|██████| 831/1865 [07:01<19:38, 1.14s/it]

```
{'eval_loss': 0.06432186812162399, 'eval_runtime': 3.0931, 'eval_samples_per_second': 120.266, 'eval_steps_per_second': 15.195, 'epoch': 2.23}
```

45%|██████| 841/1865 [07:06<19:30, 1.14s/it]

```
{'eval_loss': 0.06439536809921265, 'eval_runtime': 3.103, 'eval_samples_per_second': 119.885, 'eval_steps_per_second': 15.147, 'epoch': 2.25}
```

46%|██████| 851/1865 [07:11<19:18, 1.14s/it]

```
{'eval_loss': 0.06362882256507874, 'eval_runtime': 3.0985, 'eval_samples_per_second': 120.06, 'eval_steps_per_second': 15.169, 'epoch': 2.28}
```

46%|██████| 861/1865 [07:16<19:09, 1.14s/it]

```
{'eval_loss': 0.06394384056329727, 'eval_runtime': 3.1175, 'eval_samples_per_second': 119.326, 'eval_steps_per_second': 15.076, 'epoch': 2.31}
```

47%|██████| 871/1865 [07:21<19:06, 1.15s/it]

```
{'eval_loss': 0.06408035010099411, 'eval_runtime': 3.1478, 'eval_samples_per_second': 118.178, 'eval_steps_per_second': 14.931, 'epoch': 2.33}
```

47%|██████| 880/1865 [07:26<03:38, 4.51it/s]

```
{'eval_loss': 0.0635531513690948, 'eval_runtime': 3.1621, 'eval_samples_per_second': 117.644, 'eval_steps_per_second': 14.864, 'epoch': 2.36}
```

48%|██████| 891/1865 [07:31<18:46, 1.16s/it]

```
{'eval_loss': 0.06349756568670273, 'eval_runtime': 3.1363, 'eval_samples_per_second': 118.611, 'eval_steps_per_second': 14.986, 'epoch': 2.39}
```

48%|██████| 900/1865 [07:36<03:29, 4.61it/s]

```
{'eval_loss': 0.06337680667638779, 'eval_runtime': 3.1118, 'eval_samples_per_second': 119.545, 'eval_steps_per_second': 15.104, 'epoch': 2.41}
```

49%|██████| 911/1865 [07:41<18:29, 1.16s/it]

```
{'eval_loss': 0.06334005296230316, 'eval_runtime': 3.1486, 'eval_samples_per_second': 118.146, 'eval_steps_per_second': 14.927, 'epoch': 2.44}
```

49%|███████ | 921/1865 [07:46<18:11, 1.16s/it]

{'eval_loss': 0.0629725307226181, 'eval_runtime': 3.1474, 'eval_samples_per_second': 118.194, 'eval_steps_per_second': 14.933, 'epoch': 2.47}

50%|███████ | 931/1865 [07:51<17:55, 1.15s/it]

{'eval_loss': 0.06289377808570862, 'eval_runtime': 3.1327, 'eval_samples_per_second': 118.747, 'eval_steps_per_second': 15.003, 'epoch': 2.49}

50%|███████ | 941/1865 [07:56<17:48, 1.16s/it]

{'eval_loss': 0.0631667897105217, 'eval_runtime': 3.1346, 'eval_samples_per_second': 118.675, 'eval_steps_per_second': 14.994, 'epoch': 2.52}

51%|███████ | 951/1865 [08:01<17:35, 1.15s/it]

{'eval_loss': 0.06275201588869095, 'eval_runtime': 3.1367, 'eval_samples_per_second': 118.597, 'eval_steps_per_second': 14.984, 'epoch': 2.55}

52%|███████ | 961/1865 [08:06<17:25, 1.16s/it]

{'eval_loss': 0.06255775690078735, 'eval_runtime': 3.1397, 'eval_samples_per_second': 118.481, 'eval_steps_per_second': 14.969, 'epoch': 2.57}

52%|███████ | 971/1865 [08:11<17:21, 1.16s/it]

{'eval_loss': 0.06283076852560043, 'eval_runtime': 3.1586, 'eval_samples_per_second': 117.774, 'eval_steps_per_second': 14.88, 'epoch': 2.6}

53%|███████ | 980/1865 [08:16<03:18, 4.46it/s]

{'eval_loss': 0.06289902329444885, 'eval_runtime': 3.1331, 'eval_samples_per_second': 118.733, 'eval_steps_per_second': 15.001, 'epoch': 2.63}

53%|███████ | 991/1865 [08:21<16:53, 1.16s/it]

{'eval_loss': 0.06371808052062988, 'eval_runtime': 3.1439, 'eval_samples_per_second': 118.324, 'eval_steps_per_second': 14.95, 'epoch': 2.65}

54%|███████ | 1000/1865 [08:23<03:15, 4.43it/s]

{'loss': 0.0888, 'grad_norm': 0.06895048171281815, 'learning_rate': 0.00046380697050938335, 'epoch': 2.68}

54%|███████ | 1000/1865 [08:26<03:15, 4.43it/s]

{'eval_loss': 0.06247112154960632, 'eval_runtime': 3.1341, 'eval_samples_per_second': 118.695, 'eval_steps_per_second': 14.996, 'epoch': 2.68}

54%|███████ | 1011/1865 [08:32<16:33, 1.16s/it]

{'eval_loss': 0.06292002648115158, 'eval_runtime': 3.1485, 'eval_samples_per_second': 118.152, 'eval_steps_per_second': 14.928, 'epoch': 2.71}

55%|███████ | 1020/1865 [08:36<03:11, 4.42it/s]

{'eval_loss': 0.06193558871746063, 'eval_runtime': 3.1539, 'eval_samples_per_second': 117.95, 'eval_steps_per_second': 14.902, 'epoch': 2.73}

55%|███████ | 1031/1865 [08:42<16:03, 1.16s/it]

{'eval_loss': 0.06236349046230316, 'eval_runtime': 3.1324, 'eval_samples_per_second': 118.757, 'eval_steps_per_second': 15.004, 'epoch': 2.76}

56%|███████ | 1041/1865 [08:47<15:56, 1.16s/it]

{'eval_loss': 0.062022220343351364, 'eval_runtime': 3.1458, 'eval_samples_per_second': 118.254, 'eval_steps_per_second': 14.941, 'epoch': 2.79}

56%|███████ | 1051/1865 [08:52<15:46, 1.16s/it]

{'eval_loss': 0.06223485618829727, 'eval_runtime': 3.1529, 'eval_samples_per_second': 117.986, 'eval_steps_per_second': 14.907, 'epoch': 2.82}

57%|███████ | 1061/1865 [08:57<15:33, 1.16s/it]

{'eval_loss': 0.06181482970714569, 'eval_runtime': 3.1482, 'eval_samples_per_second': 118.163, 'eval_steps_per_second': 14.929, 'epoch': 2.84}

57%|███████ | 1071/1865 [09:02<15:20, 1.16s/it]

{'eval_loss': 0.06191721186041832, 'eval_runtime': 3.1405, 'eval_samples_per_second': 118.453, 'eval_steps_per_second': 14.966, 'epoch': 2.87}

58%|███████ | 1081/1865 [09:07<15:10, 1.16s/it]

{'eval_loss': 0.061615318059921265, 'eval_runtime': 3.1487, 'eval_samples_per_second': 118.146, 'eval_steps_per_second': 14.927, 'epoch': 2.9}

58%|███████ | 1091/1865 [09:12<15:00, 1.16s/it]

{'eval_loss': 0.06182270497083664, 'eval_runtime': 3.1543, 'eval_samples_per_second': 117.933, 'eval_steps_per_second': 14.9, 'epoch': 2.92}

59%|███████ | 1100/1865 [09:17<02:53, 4.40it/s]

```
{'eval_loss': 0.06146305799484253, 'eval_runtime': 3.1463, 'eval_samples_per_second': 118.234, 'eval_steps_per_second': 14.938, 'epoch': 2.95}
```

```
60%|███████| 1111/1865 [09:22<14:36, 1.16s/it]
```

```
{'eval_loss': 0.0617518275976181, 'eval_runtime': 3.143, 'eval_samples_per_second': 118.357, 'eval_steps_per_second': 14.954, 'epoch': 2.98}
```

```
60%|███████| 1121/1865 [09:27<14:16, 1.15s/it]
```

```
{'eval_loss': 0.061581190675497055, 'eval_runtime': 3.1535, 'eval_samples_per_second': 117.963, 'eval_steps_per_second': 14.904, 'epoch': 3.0}
```

```
61%|███████| 1131/1865 [09:32<14:16, 1.17s/it]
```

```
{'eval_loss': 0.061342302709817886, 'eval_runtime': 3.1654, 'eval_samples_per_second': 117.521, 'eval_steps_per_second': 14.848, 'epoch': 3.03}
```

```
61%|███████| 1141/1865 [09:37<13:57, 1.16s/it]
```

```
{'eval_loss': 0.06129767373204231, 'eval_runtime': 3.1319, 'eval_samples_per_second': 118.779, 'eval_steps_per_second': 15.007, 'epoch': 3.06}
```

```
62%|███████| 1151/1865 [09:42<13:49, 1.16s/it]
```

```
{'eval_loss': 0.06148406118154526, 'eval_runtime': 3.1509, 'eval_samples_per_second': 118.06, 'eval_steps_per_second': 14.916, 'epoch': 3.08}
```

```
62%|███████| 1161/1865 [09:47<13:39, 1.16s/it]
```

```
{'eval_loss': 0.06157331541180611, 'eval_runtime': 3.1553, 'eval_samples_per_second': 117.899, 'eval_steps_per_second': 14.896, 'epoch': 3.11}
```

```
63%|███████| 1171/1865 [09:52<13:30, 1.17s/it]
```

```
{'eval_loss': 0.06095639988780022, 'eval_runtime': 3.1642, 'eval_samples_per_second': 117.565, 'eval_steps_per_second': 14.854, 'epoch': 3.14}
```

```
63%|███████| 1181/1865 [09:57<13:14, 1.16s/it]
```

```
{'eval_loss': 0.06116903945803642, 'eval_runtime': 3.1497, 'eval_samples_per_second': 118.106, 'eval_steps_per_second': 14.922, 'epoch': 3.16}
```

```
64%|███████| 1191/1865 [10:02<13:04, 1.16s/it]
```

```
{'eval_loss': 0.06092489883303642, 'eval_runtime': 3.1533, 'eval_samples_per_second': 117.971, 'eval_steps_per_second': 14.905, 'epoch': 3.19}
```

64%|███████ | 1200/1865 [10:07<02:31, 4.39it/s]
{'eval_loss': 0.060990527272224426, 'eval_runtime': 3.1583, 'eval_samples_per_second': 117.785, 'eval_steps_per_second': 14.881, 'epoch': 3.22}

65%|███████ | 1211/1865 [10:13<12:43, 1.17s/it]
{'eval_loss': 0.06089864671230316, 'eval_runtime': 3.1593, 'eval_samples_per_second': 117.749, 'eval_steps_per_second': 14.877, 'epoch': 3.24}

65%|███████ | 1221/1865 [10:18<12:27, 1.16s/it]
{'eval_loss': 0.060612503439188004, 'eval_runtime': 3.1453, 'eval_samples_per_second': 118.271, 'eval_steps_per_second': 14.943, 'epoch': 3.27}

66%|███████ | 1231/1865 [10:23<12:21, 1.17s/it]
{'eval_loss': 0.0606440044939518, 'eval_runtime': 3.1714, 'eval_samples_per_second': 117.298, 'eval_steps_per_second': 14.82, 'epoch': 3.3}

67%|███████ | 1241/1865 [10:28<12:07, 1.17s/it]
{'eval_loss': 0.06070176139473915, 'eval_runtime': 3.1617, 'eval_samples_per_second': 117.659, 'eval_steps_per_second': 14.866, 'epoch': 3.32}

67%|███████ | 1251/1865 [10:33<11:56, 1.17s/it]
{'eval_loss': 0.06057312712073326, 'eval_runtime': 3.1614, 'eval_samples_per_second': 117.668, 'eval_steps_per_second': 14.867, 'epoch': 3.35}

68%|███████ | 1261/1865 [10:38<11:44, 1.17s/it]
{'eval_loss': 0.060515373945236206, 'eval_runtime': 3.1634, 'eval_samples_per_second': 117.593, 'eval_steps_per_second': 14.857, 'epoch': 3.38}

68%|███████ | 1271/1865 [10:43<11:32, 1.17s/it]
{'eval_loss': 0.060686007142066956, 'eval_runtime': 3.1601, 'eval_samples_per_second': 117.719, 'eval_steps_per_second': 14.873, 'epoch': 3.4}

69%|███████ | 1281/1865 [10:48<11:21, 1.17s/it]
{'eval_loss': 0.06065713241696358, 'eval_runtime': 3.1646, 'eval_samples_per_second': 117.549, 'eval_steps_per_second': 14.852, 'epoch': 3.43}

69%|███████ | 1291/1865 [10:53<11:09, 1.17s/it]

```
{'eval_loss': 0.06052849814295769, 'eval_runtime': 3.1598, 'eval_samples_per_second': 117.728, 'eval_steps_per_second': 14.874, 'epoch': 3.46}
```

70%|███████ | 1300/1865 [10:58<02:09, 4.37it/s]

```
{'eval_loss': 0.06049437075853348, 'eval_runtime': 3.153, 'eval_samples_per_second': 117.984, 'eval_steps_per_second': 14.907, 'epoch': 3.49}
```

70%|███████ | 1311/1865 [11:04<10:45, 1.17s/it]

```
{'eval_loss': 0.06041036546230316, 'eval_runtime': 3.1549, 'eval_samples_per_second': 117.912, 'eval_steps_per_second': 14.898, 'epoch': 3.51}
```

71%|███████ | 1321/1865 [11:09<10:33, 1.17s/it]

```
{'eval_loss': 0.06034473702311516, 'eval_runtime': 3.1586, 'eval_samples_per_second': 117.775, 'eval_steps_per_second': 14.88, 'epoch': 3.54}
```

71%|███████ | 1331/1865 [11:14<10:25, 1.17s/it]

```
{'eval_loss': 0.06049174815416336, 'eval_runtime': 3.1793, 'eval_samples_per_second': 117.006, 'eval_steps_per_second': 14.783, 'epoch': 3.57}
```

72%|███████ | 1341/1865 [11:19<10:09, 1.16s/it]

```
{'eval_loss': 0.06027385592460632, 'eval_runtime': 3.1525, 'eval_samples_per_second': 118.003, 'eval_steps_per_second': 14.909, 'epoch': 3.59}
```

72%|███████ | 1351/1865 [11:24<09:57, 1.16s/it]

```
{'eval_loss': 0.060386739671230316, 'eval_runtime': 3.1509, 'eval_samples_per_second': 118.062, 'eval_steps_per_second': 14.916, 'epoch': 3.62}
```

73%|███████ | 1361/1865 [11:29<09:45, 1.16s/it]

```
{'eval_loss': 0.06023447960615158, 'eval_runtime': 3.1534, 'eval_samples_per_second': 117.967, 'eval_steps_per_second': 14.904, 'epoch': 3.65}
```

74%|███████ | 1371/1865 [11:34<09:57, 1.21s/it]

```
{'eval_loss': 0.06025548279285431, 'eval_runtime': 3.1619, 'eval_samples_per_second': 117.651, 'eval_steps_per_second': 14.865, 'epoch': 3.67}
```

74%|███████ | 1381/1865 [11:39<09:24, 1.17s/it]

```
{'eval_loss': 0.06031585857272148, 'eval_runtime': 3.1625, 'eval_samples_per_second': 117.627, 'eval_steps_per_second': 14.861, 'epoch': 3.7}
```


75%|███████ | 1391/1865 [11:44<09:11, 1.16s/it]
{'eval_loss': 0.0602213554084301, 'eval_runtime': 3.1567, 'eval_samples_per_second': 117.846, 'eval_steps_per_second': 14.889, 'epoch': 3.73}

75%|███████ | 1400/1865 [11:49<01:45, 4.39it/s]
{'eval_loss': 0.060142599046230316, 'eval_runtime': 3.1484, 'eval_samples_per_second': 118.154, 'eval_steps_per_second': 14.928, 'epoch': 3.75}

76%|███████ | 1411/1865 [11:55<08:50, 1.17s/it]
{'eval_loss': 0.05994571000337601, 'eval_runtime': 3.1597, 'eval_samples_per_second': 117.731, 'eval_steps_per_second': 14.875, 'epoch': 3.78}

76%|███████ | 1421/1865 [12:00<08:38, 1.17s/it]
{'eval_loss': 0.06015310063958168, 'eval_runtime': 3.1601, 'eval_samples_per_second': 117.717, 'eval_steps_per_second': 14.873, 'epoch': 3.81}

77%|███████ | 1431/1865 [12:05<08:25, 1.16s/it]
{'eval_loss': 0.06013735011219978, 'eval_runtime': 3.1552, 'eval_samples_per_second': 117.901, 'eval_steps_per_second': 14.896, 'epoch': 3.83}

77%|███████ | 1440/1865 [12:10<01:36, 4.40it/s]
{'eval_loss': 0.05998246371746063, 'eval_runtime': 3.1542, 'eval_samples_per_second': 117.937, 'eval_steps_per_second': 14.901, 'epoch': 3.86}

78%|███████ | 1451/1865 [12:15<08:00, 1.16s/it]
{'eval_loss': 0.059974588453769684, 'eval_runtime': 3.1464, 'eval_samples_per_second': 118.231, 'eval_steps_per_second': 14.938, 'epoch': 3.89}

78%|███████ | 1461/1865 [12:20<07:50, 1.17s/it]
{'eval_loss': 0.05992733687162399, 'eval_runtime': 3.1653, 'eval_samples_per_second': 117.526, 'eval_steps_per_second': 14.849, 'epoch': 3.91}

79%|███████ | 1471/1865 [12:25<07:40, 1.17s/it]
{'eval_loss': 0.05984333157539368, 'eval_runtime': 3.1685, 'eval_samples_per_second': 117.405, 'eval_steps_per_second': 14.833, 'epoch': 3.94}

79%|███████ | 1481/1865 [12:30<07:28, 1.17s/it]

```
{'eval_loss': 0.0599115826189518, 'eval_runtime': 3.1615, 'eval_samples_per_second': 117.666, 'eval_steps_per_second': 14.866, 'epoch': 3.97}
```

```
80%|██████████ | 1491/1865 [12:35<07:14, 1.16s/it]
```

```
{'eval_loss': 0.05967269465327263, 'eval_runtime': 3.1466, 'eval_samples_per_second': 118.224, 'eval_steps_per_second': 14.937, 'epoch': 3.99}
```

```
80%|██████████ | 1500/1865 [12:37<01:27, 4.16it/s]
```

```
{'loss': 0.0785, 'grad_norm': 0.04232126846909523, 'learning_rate': 0.00019571045576407506, 'epoch': 4.02}
```

```
80%|██████████ | 1500/1865 [12:40<01:27, 4.16it/s]
```

```
{'eval_loss': 0.05969894677400589, 'eval_runtime': 3.1481, 'eval_samples_per_second': 118.165, 'eval_steps_per_second': 14.929, 'epoch': 4.02}
```

```
81%|██████████ | 1510/1865 [12:45<01:22, 4.32it/s]
```

```
{'eval_loss': 0.05979344993829727, 'eval_runtime': 3.1578, 'eval_samples_per_second': 117.804, 'eval_steps_per_second': 14.884, 'epoch': 4.05}
```

```
82%|██████████ | 1521/1865 [12:50<06:41, 1.17s/it]
```

```
{'eval_loss': 0.05973832309246063, 'eval_runtime': 3.163, 'eval_samples_per_second': 117.61, 'eval_steps_per_second': 14.859, 'epoch': 4.08}
```

```
82%|██████████ | 1531/1865 [12:55<06:29, 1.16s/it]
```

```
{'eval_loss': 0.05989058315753937, 'eval_runtime': 3.1556, 'eval_samples_per_second': 117.884, 'eval_steps_per_second': 14.894, 'epoch': 4.1}
```

```
83%|██████████ | 1540/1865 [13:00<01:14, 4.38it/s]
```

```
{'eval_loss': 0.059838078916072845, 'eval_runtime': 3.1586, 'eval_samples_per_second': 117.772, 'eval_steps_per_second': 14.88, 'epoch': 4.13}
```

```
83%|██████████ | 1551/1865 [13:06<06:06, 1.17s/it]
```

```
{'eval_loss': 0.059572938829660416, 'eval_runtime': 3.1631, 'eval_samples_per_second': 117.605, 'eval_steps_per_second': 14.859, 'epoch': 4.16}
```

```
84%|██████████ | 1560/1865 [13:10<01:09, 4.40it/s]
```

```
{'eval_loss': 0.05951518565416336, 'eval_runtime': 3.1498, 'eval_samples_per_second': 118.103, 'eval_steps_per_second': 14.922, 'epoch': 4.18}
```

```
84%|██████████ | 1571/1865 [13:16<05:42, 1.17s/it]
```

```
{'eval_loss': 0.05953618511557579, 'eval_runtime': 3.1603, 'eval_samples_per_second': 117.709, 'eval_steps_per_second': 14.872, 'epoch': 4.21}
```

85%|██████████ | 1581/1865 [13:21<05:30, 1.16s/it]

```
{'eval_loss': 0.059546686708927155, 'eval_runtime': 3.1552, 'eval_samples_per_second': 117.9, 'eval_steps_per_second': 14.896, 'epoch': 4.24}
```

85%|██████████ | 1591/1865 [13:26<05:19, 1.16s/it]

```
{'eval_loss': 0.05962543934583664, 'eval_runtime': 3.1591, 'eval_samples_per_second': 117.757, 'eval_steps_per_second': 14.878, 'epoch': 4.26}
```

86%|██████████ | 1600/1865 [13:31<01:00, 4.37it/s]

```
{'eval_loss': 0.05946793034672737, 'eval_runtime': 3.1572, 'eval_samples_per_second': 117.824, 'eval_steps_per_second': 14.886, 'epoch': 4.29}
```

86%|██████████ | 1611/1865 [13:36<04:56, 1.17s/it]

```
{'eval_loss': 0.059328798204660416, 'eval_runtime': 3.155, 'eval_samples_per_second': 117.906, 'eval_steps_per_second': 14.897, 'epoch': 4.32}
```

87%|██████████ | 1620/1865 [13:41<00:55, 4.38it/s]

```
{'eval_loss': 0.059478431940078735, 'eval_runtime': 3.1527, 'eval_samples_per_second': 117.996, 'eval_steps_per_second': 14.908, 'epoch': 4.34}
```

87%|██████████ | 1631/1865 [13:46<04:33, 1.17s/it]

```
{'eval_loss': 0.059633318334817886, 'eval_runtime': 3.1661, 'eval_samples_per_second': 117.493, 'eval_steps_per_second': 14.845, 'epoch': 4.37}
```

88%|██████████ | 1641/1865 [13:51<04:20, 1.16s/it]

```
{'eval_loss': 0.05948368087410927, 'eval_runtime': 3.1519, 'eval_samples_per_second': 118.025, 'eval_steps_per_second': 14.912, 'epoch': 4.4}
```

89%|██████████ | 1651/1865 [13:56<04:07, 1.16s/it]

```
{'eval_loss': 0.059504684060811996, 'eval_runtime': 3.1386, 'eval_samples_per_second': 118.523, 'eval_steps_per_second': 14.975, 'epoch': 4.42}
```

89%|██████████ | 1661/1865 [14:01<03:57, 1.16s/it]

```
{'eval_loss': 0.05947580561041832, 'eval_runtime': 3.1523, 'eval_samples_per_second': 118.008, 'eval_steps_per_second': 14.91, 'epoch': 4.45}
```

```
90%|██████████ | 1671/1865 [14:06<03:45, 1.16s/it]
{'eval_loss': 0.05935505032539368, 'eval_runtime': 3.1525, 'eval_samples_per_second': 118.002, 'eval_steps_per_second': 14.909, 'epoch': 4.48}

90%|██████████ | 1681/1865 [14:12<03:34, 1.17s/it]
{'eval_loss': 0.059234291315078735, 'eval_runtime': 3.1612, 'eval_samples_per_second': 117.675, 'eval_steps_per_second': 14.868, 'epoch': 4.5}

91%|██████████ | 1691/1865 [14:17<03:22, 1.16s/it]
{'eval_loss': 0.05947580561041832, 'eval_runtime': 3.1528, 'eval_samples_per_second': 117.989, 'eval_steps_per_second': 14.907, 'epoch': 4.53}

91%|██████████ | 1700/1865 [14:21<00:37, 4.40it/s]
{'eval_loss': 0.05953618511557579, 'eval_runtime': 3.151, 'eval_samples_per_second': 118.057, 'eval_steps_per_second': 14.916, 'epoch': 4.56}

92%|██████████ | 1711/1865 [14:27<02:59, 1.16s/it]
{'eval_loss': 0.05943642929196358, 'eval_runtime': 3.16, 'eval_samples_per_second': 117.72, 'eval_steps_per_second': 14.873, 'epoch': 4.58}

92%|██████████ | 1721/1865 [14:32<02:48, 1.17s/it]
{'eval_loss': 0.05929729342460632, 'eval_runtime': 3.1699, 'eval_samples_per_second': 117.355, 'eval_steps_per_second': 14.827, 'epoch': 4.61}

93%|██████████ | 1731/1865 [14:37<02:36, 1.17s/it]
{'eval_loss': 0.05927629396319389, 'eval_runtime': 3.1641, 'eval_samples_per_second': 117.571, 'eval_steps_per_second': 14.854, 'epoch': 4.64}

93%|██████████ | 1741/1865 [14:42<02:24, 1.17s/it]
{'eval_loss': 0.05921328812837601, 'eval_runtime': 3.1629, 'eval_samples_per_second': 117.613, 'eval_steps_per_second': 14.86, 'epoch': 4.66}

94%|██████████ | 1751/1865 [14:47<02:12, 1.16s/it]
{'eval_loss': 0.05922378972172737, 'eval_runtime': 3.1465, 'eval_samples_per_second': 118.227, 'eval_steps_per_second': 14.937, 'epoch': 4.69}

94%|██████████ | 1761/1865 [14:52<02:01, 1.17s/it]
```

```
{'eval_loss': 0.059315670281648636, 'eval_runtime': 3.1629, 'eval_samples_per_second': 117.615, 'eval_steps_per_second': 14.86, 'epoch': 4.72}
```

```
95%|██████████ | 1771/1865 [14:57<01:49, 1.16s/it]
```

```
{'eval_loss': 0.05932092294096947, 'eval_runtime': 3.1534, 'eval_samples_per_second': 117.968, 'eval_steps_per_second': 14.905, 'epoch': 4.75}
```

```
95%|██████████ | 1780/1865 [15:02<00:19, 4.38it/s]
```

```
{'eval_loss': 0.05931304767727852, 'eval_runtime': 3.1666, 'eval_samples_per_second': 117.477, 'eval_steps_per_second': 14.842, 'epoch': 4.77}
```

```
96%|██████████ | 1791/1865 [15:07<01:26, 1.17s/it]
```

```
{'eval_loss': 0.05930516868829727, 'eval_runtime': 3.1578, 'eval_samples_per_second': 117.803, 'eval_steps_per_second': 14.884, 'epoch': 4.8}
```

```
97%|██████████ | 1800/1865 [15:12<00:14, 4.39it/s]
```

```
{'eval_loss': 0.059389177709817886, 'eval_runtime': 3.1528, 'eval_samples_per_second': 117.989, 'eval_steps_per_second': 14.907, 'epoch': 4.83}
```

```
97%|██████████ | 1811/1865 [15:18<01:02, 1.17s/it]
```

```
{'eval_loss': 0.059315670281648636, 'eval_runtime': 3.1538, 'eval_samples_per_second': 117.951, 'eval_steps_per_second': 14.902, 'epoch': 4.85}
```

```
98%|██████████ | 1821/1865 [15:23<00:51, 1.16s/it]
```

```
{'eval_loss': 0.05927892029285431, 'eval_runtime': 3.1529, 'eval_samples_per_second': 117.985, 'eval_steps_per_second': 14.907, 'epoch': 4.88}
```

```
98%|██████████ | 1831/1865 [15:28<00:39, 1.16s/it]
```

```
{'eval_loss': 0.05930516868829727, 'eval_runtime': 3.1419, 'eval_samples_per_second': 118.401, 'eval_steps_per_second': 14.959, 'epoch': 4.91}
```

```
99%|██████████ | 1841/1865 [15:33<00:27, 1.16s/it]
```

```
{'eval_loss': 0.059289418160915375, 'eval_runtime': 3.1462, 'eval_samples_per_second': 118.238, 'eval_steps_per_second': 14.939, 'epoch': 4.93}
```

```
99%|██████████ | 1851/1865 [15:38<00:16, 1.17s/it]
```

```
{'eval_loss': 0.05931304767727852, 'eval_runtime': 3.1648, 'eval_samples_per_second': 117.542, 'eval_steps_per_second': 14.851, 'epoch': 4.96}
```

```
100%|██████████| 1861/1865 [15:43<00:04, 1.16s/it]
{'eval_loss': 0.05931304767727852, 'eval_runtime': 3.1507, 'eval_samples_per_second': 118.07, 'eval_steps_per_second': 14.917, 'epoch': 4.99}

100%|██████████| 1865/1865 [15:44<00:00, 1.97it/s]
{'train_runtime': 944.3468, 'train_samples_per_second': 15.778, 'train_steps_per_second': 1.975, 'train_loss': 0.4984449616705764, 'epoch': 5.0}
```

```
In [ ]: # Evaluate using the validation dataset
peft_trainer.evaluate(tokenized_datasets['validation'])
```

```
100%|██████████| 47/47 [00:03<00:00, 15.16it/s]
```

```
Out[ ]: {'eval_loss': 0.05931829661130905,
        'eval_runtime': 3.1482,
        'eval_samples_per_second': 118.162,
        'eval_steps_per_second': 14.929,
        'epoch': 5.0}
```

```
In [ ]: # Path to the log file
import json
log_file_path = f"{output_dir}/checkpoint-1865/trainer_state.json"

# Load the log file
with open(log_file_path, "r") as log_file:
    logs = json.load(log_file)

# Extract the training losses
training_losses = [entry["loss"] for entry in logs["log_history"] if "loss" in entry]
# Extract validation loss from log history
validation_loss = [entry['eval_loss'] for entry in logs["log_history"] if "eval_loss" in entry]

# Plot the learning curve
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))

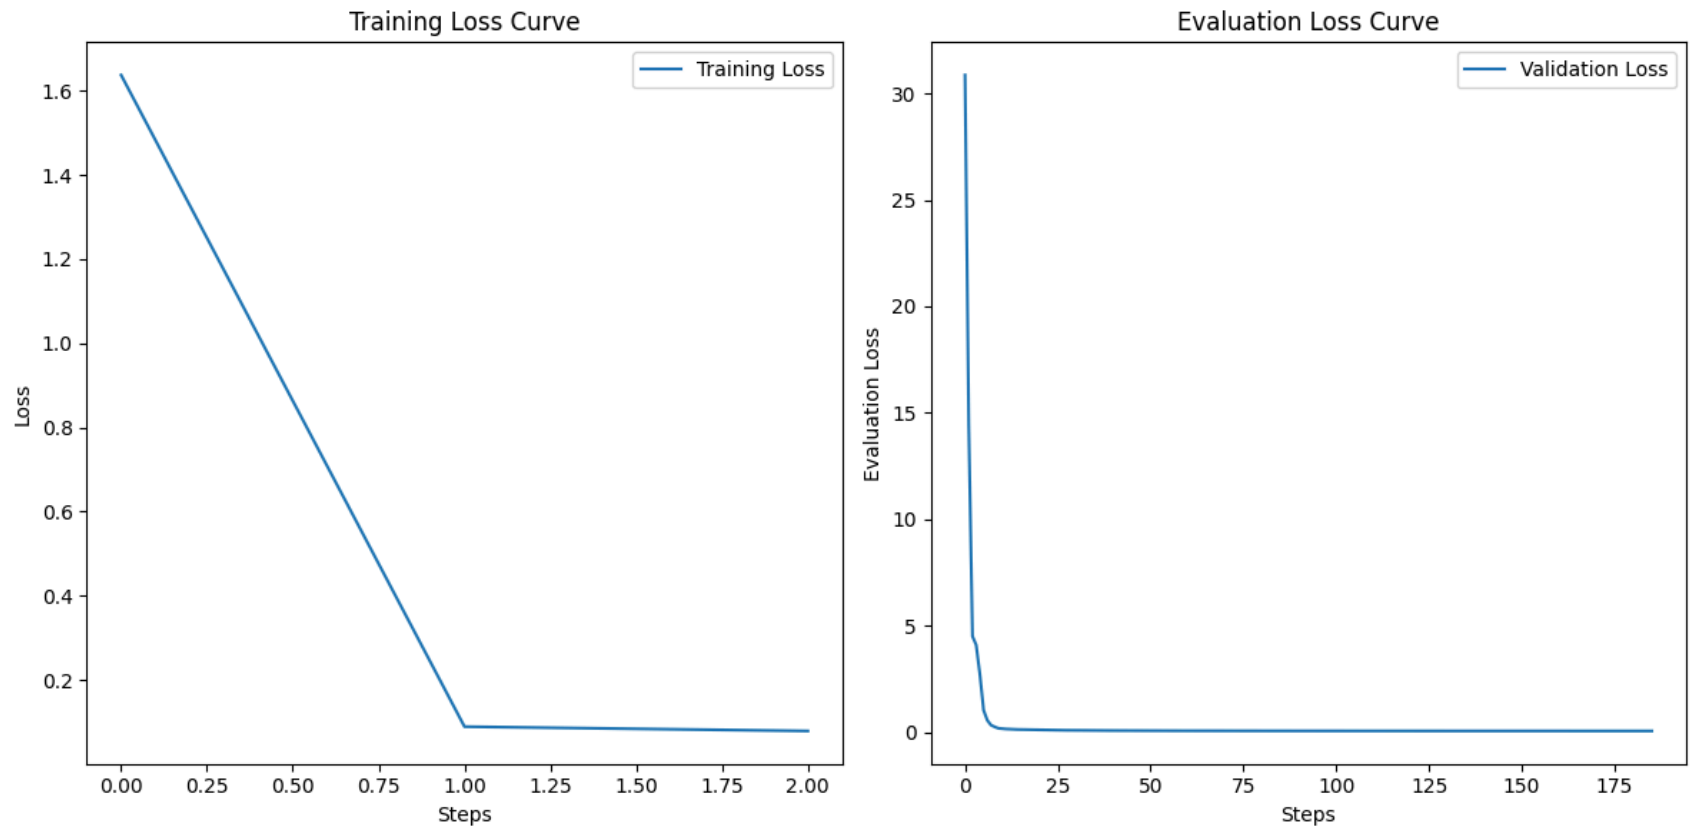
ax1.plot(training_losses, label="Training Loss")
ax1.set_xlabel("Steps")
ax1.set_ylabel("Loss")
ax1.set_title("Training Loss Curve")
ax1.legend()

ax2.plot(validation_loss, label="Validation Loss")
```

```
ax2.set_xlabel("Steps")
ax2.set_ylabel("Evaluation Loss")
ax2.set_title("Evaluation Loss Curve")
ax2.legend()

plt.tight_layout()

plt.show()
```



Save the fine-tuned model

```
In [ ]: peft_model_path="./peft-conversation-checkpoint-local"

peft_trainer.model.save_pretrained(peft_model_path)
tokenizer.save_pretrained(peft_model_path)
```

```
Out[ ]: ('./peft-conversation-checkpoint-local\\tokenizer_config.json',
        './peft-conversation-checkpoint-local\\special_tokens_map.json',
        './peft-conversation-checkpoint-local\\tokenizer.json')
```

Adding an adapter to the original flan-t5 model

```
In [ ]: from peft import PeftModel
```

```
peft_model_base = AutoModelForSeq2SeqLM.from_pretrained("google/flan-t5-base", torch_dtype=torch.bfloat16)
tokenizer = AutoTokenizer.from_pretrained("google/flan-t5-base")

peft_model = PeftModel.from_pretrained(peft_model_base,
                                       'peft-conversation-checkpoint-local',
                                       torch_dtype=torch.bfloat16,
                                       is_trainable=False)
```

```
In [ ]: print(model_parameters(peft_model))
```

```
trainable model parameters: 0
all model parameters: 248462592
percentage of trainable model parameters: 0.00%
```

```
In [ ]: def insert_prompt(prompt):
        input_ids = tokenizer(prompt, return_tensors="pt").input_ids

        peft_model_outputs = peft_model.generate(input_ids=input_ids, generation_config=GenerationConfig(max_new_tokens=128))
        peft_model_text_output = tokenizer.decode(peft_model_outputs[0], skip_special_tokens=True)

        return peft_model_text_output
```

```
In [ ]: prompt = "where is paris?"

print(f'question:\n{prompt}')
print(f'PEFT MODEL: {insert_prompt(prompt)}')
```

```
question:
where is paris?
PEFT MODEL: France
```

```
In [ ]: prompt = "What is the capital of california?"
```



```
print(f'question:\n{prompt}')
```

```
print(f'PEFT MODEL: {insert_prompt(prompt)}')
```

question:

What is the capital of california?

PEFT MODEL: San Francisco

```
In [ ]: prompt = "What is the capital of Italy"
```

```
print(f'question:\n{prompt}')
```

```
print(f'PEFT MODEL: {insert_prompt(prompt)}')
```

question:

What is the capital of Italy

PEFT MODEL: Rome

```
In [ ]: prompt = "What is the best drink for a hangover?"
```

```
print(f'question:\n{prompt}')
```

```
print(f'PEFT MODEL: {insert_prompt(prompt)}')
```

question:

What is the best drink for a hangover?

PEFT MODEL: a stout

```
In [ ]: prompt = "What is the probability that two apples will fall at the same time?"
```

```
print(f'question:\n{prompt}')
```

```
print(f'PEFT MODEL: {insert_prompt(prompt)}')
```

question:

What is the probability that two apples will fall at the same time?

PEFT MODEL: 1 / 2